

simplicio-cli benchmark

Benchmark — simplicio-cli (offline harness)

Date: **2026-05-26**

Models: qwen/qwen-2.5-7b-instruct, meta-llama/llama-3.1-8b-instruct, google/gemma-3-12b-it


Cases: **10** across stacks: angular, dotnet, react

Base: <https://openrouter.ai/api/v1>


Each check is a deterministic regex against the model output (target-file mention, DIFF block, TEST block, contract-state words). Same model on both sides — only the prompt structure changes. The *without* run is the raw one-line objetivo; the *with* run wraps the same objetivo in simplicio's 6-layer contract.

Headline

- **Without simplicio:** 58/156 (37%)
- **With simplicio:** 142/156 (91%)
- **Delta: +54 points** (+145% relative)




pass rate by model



Per-model breakdown

Model	Cases	Without	With	Delta (pts)	Relative gain
qwen/qwen-2.5-7b-instruct	10	20/52 (38%)	42/52 (80%)	+42	+110%
meta-llama/llama-3.1-8b-instruct	10	18/52 (34%)	51/52 (98%)	+64	+183%
google/gemma-3-12b-it	10	20/52 (38%)	49/52 (94%)	+56	+145%

Per-case (averaged across models)




per case

#	Stack	Task	Without	With	Δ
1	angular	Hide the Delete button when the current user is not an admin	40%	73%	+33
2	angular	Disable the email field unless the profile role is editor.	40%	80%	+40
3	angular	Only show the audit log link for users with role 'auditor'.	27%	100%	+73
4	angular	Show 'Approve' button only when the order status is 'pending	39%	100%	+61
5	react	Render the export menu item only for users in the 'analytics	33%	100%	+67
6	react	Disable the 'Save Draft' button while the form is invalid OR	50%	100%	+50
7	react	Show a 'No results' empty state when the search returns zero	27%	100%	+73
8	dotnet	Require the 'CanApprove' policy on the Approve endpoint of t	47%	80%	+33
9	dotnet	Restrict the GET /reports endpoint so only users in the Mana	27%	73%	+47
10	angular	Show a warning banner if the	40%	100%	+60

#	Stack	Task	Without	With	Δ
		user has unsaved changes and tr			

Per-stack



per stack

Stack	Without	With	Δ
angular	37%	91%	+54
dotnet	37%	77%	+40
react	38%	100%	+62

Output-quality signals (rate across all runs)

Beyond pass-rate, the same outputs are scored on structural quality. Each row = % of runs (cases $\times$ models) where the signal is present.

Signal	Without simplicio	With simplicio
DIFF block present	0% (0/30)	100% (30/30)
TEST block present	86% (26/30)	93% (28/30)
target file mentioned	3% (1/30)	96% (29/30)
avg criteria-keywords hit / run	9.4	9.5
avg output length (chars)	3044	2106

How to reproduce

```
OPENROUTER_API_KEY=... \  
  BENCH_MODELS="qwen/qwen-2.5-7b-instruct,meta-llama/llama-3.1-8b-  
    instruct,google/gemma-3-12b-it" \  
python3 bench/run_offline.py
```

Raw model outputs are saved under
.simplicio/bench_runs/<model>/case_NN/{sem,com}.txt so you can audit what
the LLM actually produced on each side. Charts are SVG under bench/charts/; raw
aggregated data under bench/results.json.