

simplicio-cli — bench results

Benchmark — simplicio-cli (offline harness)

Date: **2026-05-26**

Models: qwen/qwen-2.5-7b-instruct, meta-llama/llama-3.1-8b-instruct, google/gemma-3-12b-it


Cases: **10** across stacks: angular, dotnet, react

Base: <https://openrouter.ai/api/v1>


Each check is a deterministic regex against the model output (target-file mention, DIFF block, TEST block, contract-state words). Same model on both sides — only the prompt structure changes. The *without* run is the raw one-line objetivo; the *with* run wraps the same objetivo in simplicio's 6-layer contract.

Headline

- **Without simplicio:** 55/156 (35%)
- **With simplicio:** 141/156 (90%)
- **Delta: +55 points** (+156% relative)




pass rate by model



Per-model breakdown

Model	Cases	Without	With	Delta (pts)	Relative gain
qwen/qwen-2.5-7b-instruct	10	18/52 (34%)	46/52 (88%)	+54	+156%
meta-llama/llama-3.1-8b-instruct	10	19/52 (36%)	47/52 (90%)	+54	+147%
google/gemma-3-12b-it	10	18/52 (34%)	48/52 (92%)	+58	+167%

Per-case (averaged across models)




per case

#	Stack	Task	Without	With	Δ
1	angular	Hide the Delete button when the current user is not an admin	40%	73%	+33
2	angular	Disable the email field unless the profile role is editor.	40%	100%	+60
3	angular	Only show the audit log link for users with role 'auditor'.	20%	93%	+73
4	angular	Show 'Approve' button only when the order status is 'pending	39%	100%	+61
5	react	Render the export menu item only for users in the 'analytics	33%	93%	+60
6	react	Disable the 'Save Draft' button while the form is invalid OR	50%	94%	+44
7	react	Show a 'No results' empty state when the search returns zero	27%	93%	+67
8	dotnet	Require the 'CanApprove' policy on the Approve endpoint of t	33%	93%	+60
9	dotnet	Restrict the GET /reports endpoint so only users in the Mana	27%	73%	+47
10	angular	Show a warning banner if the	40%	87%	+47

#	Stack	Task	Without	With	Δ
		user has unsaved changes and tr			

Per-stack



per stack

Stack	Without	With	Δ
angular	36%	91%	+55
dotnet	30%	83%	+53
react	38%	94%	+56

Output-quality signals (rate across all runs)

Beyond pass-rate, the same outputs are scored on structural quality. Each row = % of runs (cases $\times$ models) where the signal is present.

Signal	Without simplicio	With simplicio
DIFF block present	0% (0/30)	100% (30/30)
TEST block present	80% (24/30)	96% (29/30)
target file mentioned	0% (0/30)	96% (29/30)
avg criteria-keywords hit / run	9.5	9.6
avg output length (chars)	3045	2111

Cost — tokens & wall-clock (measured, per run)

Reported straight from the provider's usage field and `time.perf_counter()`. *Per-run* = one model call (one case, one side). With simplicio uses more input tokens (the 6-layer wrap) and fewer output tokens (model stops guessing earlier).

Model	Side	Avg prompt tok	Avg completion tok	Avg total tok	Avg latency
qwen/qwen-2.5-7b-instruct	without	46	689	735	11871 ms
qwen/qwen-2.5-7b-instruct	with	228	534	763	9129 ms
meta-llama/llama-3.1-8b-instruct	without	29	586	616	11451 ms
meta-llama/llama-3.1-8b-instruct	with	213	470	683	10358 ms
google/gemma-3-12b-it	without	25	900	925	13983 ms
google/gemma-3-12b-it	with	225	640	866	10341 ms

Aggregate over the full bench (30 runs per side):

- without simplicio: 22,774 tokens total · 373.1s wall-clock · 759 tok/run · 12435 ms/run
- with simplicio: 23,127 tokens total · 298.3s wall-clock · 770 tok/run · 9943 ms/run
- token delta: +353 (+1%)
- time delta: -74.8s (-21%)

How to reproduce

```
OPENROUTER_API_KEY=... \  
  BENCH_MODELS="qwen/qwen-2.5-7b-instruct,meta-llama/llama-3.1-8b-  
    instruct,google/gemma-3-12b-it" \  
python3 bench/run_offline.py
```

Raw model outputs are saved under

`.simplicio/bench_runs/<model>/case_NN/{sem,com}.txt` so you can audit what the LLM actually produced on each side. Charts are SVG under `bench/charts/`; raw aggregated data under `bench/results.json`.