

Benchmark - old vs new (with & without simplicio-cli)

Date: 2026-05-28

Old = pass-rate published in the README. New = re-run on the latest version: Qwen2.5-Coder and HF-served models via the Hugging Face router, the rest via OpenRouter. Same 10 cases/side, deterministic regex checks. 'n/a' = not reproducible in this environment.

Local offline - qwen2.5-coder (was Ollama; now HF)

Model	w/o old	w/o new	with old	with new	Dw/o	Dw
Qwen 2.5 Coder 7B	36%	38%	92%	96%	+2	+4
Qwen 2.5 Coder 3B	34%	34%	82%	94%	+0	+12
Qwen 2.5 Coder 1.5B	32%	30%	88%	92%	-2	+4

Tiny models - sub-4B

Model	w/o old	w/o new	with old	with new	Dw/o	Dw
Gemma 3 4B	38%	40%	96%	92%	+2	-4
Llama 3.2 3B	28%	30%	73%	76%	+2	+3
Gemma 3n e4B	44%	38%	88%	90%	-6	+2
Phi-4 mini	36%	36%	73%	73%	+0	+0
Llama 3.2 1B	26%	25%	40%	36%	-1	-4

Frontier 2026 models

Model	w/o old	w/o new	with old	with new	Dw/o	Dw
GPT-5.5	38%	38%	100%	98%	+0	-2
Kimi K2.6	40%	44%	100%	100%	+4	+0
Gemini 3.5 Flash	42%	42%	100%	100%	+0	+0
Qwen 3.7 Max	44%	n/a	100%	n/a	n/a	n/a
Claude Opus 4.7	42%	n/a	98%	n/a	n/a	n/a
DeepSeek V4 Pro	44%	n/a	96%	n/a	n/a	n/a

Mid-tier 7B-12B open models

Model	w/o old	w/o new	with old	with new	Dw/o	Dw
Gemma 3 12B	34%	46%	92%	92%	+12	+0
Llama 3.1 8B	36%	36%	90%	88%	+0	-2
Qwen 2.5 7B	34%	34%	88%	100%	+0	+12

Overall (models with a new re-run)

Side	Old avg	New avg	Delta
Without simplicio	36%	36%	+1
With simplicio	86%	88%	+2

Models re-run: 14. Merged new dataset: bench/results_all.json.