

Execution benchmark - real project, real tasks, real test suite

Date: 2026-05-28 Target: wesleysimplicio/sistema-sindico (real PHP 8 condominium system, public on GitHub) Tasks: 12
additive engineering changes across Core / Middleware / Repository / Router.

Methodology - what 'pass' actually means

NOT regex on model output. NOT a toy harness in isolation. The model's generated file replaces the real source file in a working copy of the real repo; a hidden PHPUnit test (never shown to the model, asserting TRUE and FALSE states) is added to tests/unit/Core/Hidden/; the FULL production suite runs (existing tests + hidden). Pass = vendor/bin/phpunit exit code 0 - same green/red the project's CI would use to merge a PR. The model's change must be correct AND must not break any pre-existing test of the real codebase.

All sides emit the complete file (identical output shape). Variables: baseline = raw goal; cli = the simplicio-cli 6-layer task contract; sp = the simplicio-prompt Tuple-Space + Yool runtime template.

Headline

Side	Passed	Rate	vs baseline
Baseline	17/36	47%	-
simplicio-cli (6-layer)	27/36	75%	+28 pts
simplicio-prompt (Yool)	16/36	44%	-3 pts

Per-model (pass = full PHPUnit suite green)

Model	Baseline	cli	sp	D cli	D sp
google/gemma-3-4b-it	4/12 (33%)	8/12 (66%)	4/12 (33%)	+33	+0
meta-llama/llama-3.1-8b-instruct	5/12 (41%)	7/12 (58%)	4/12 (33%)	+17	-8
google/gemini-3.5-flash	8/12 (66%)	12/12 (100%)	8/12 (66%)	+34	+0

Per-task x model (base / cli / sp)

Task	gemma-3-4b-it	llama-3.1-8b-instruct	gemini-3.5-flash
password_strength	/P/.	/P/.	P/P/P
password_require_symbol	./.	P/P/.	P/P/.
env_get_int	/P/.	./.	/P/.
env_get_bool	/P/.	./.	/P/.
admin_only_allowed_roles	P/P/P	P/P/P	P/P/P
rate_limit_bucket_key	/P/.	/P/.	/P/P
base_repository_build_where_sql	./.	./.	/P/.
router_has	P/P/P	P/P/P	P/P/P
bugfix_password_policy_lower_case	P/P/P	P/P/P	P/P/P
password_assess	P/P/P	P/P/P	P/P/P
base_repository_build_update_sql	./.	./.	P/P/P
router_extract_params	./.	./.	P/P/P