

Execution benchmark - sistema-sindico (real PHPUnit, not regex)

Date: 2026-05-28 Tasks: 4 additive PHP modifications

Three prompt variants per task, all emitting the complete file. Pass = the full PHPUnit suite (existing + hidden test) goes green: the new method works AND nothing else broke. Variables: baseline = raw goal; cli = the simplicio-cli 6-layer task contract; sp = the simplicio-prompt Tuple-Space + Yool runtime template with the task as user input X.

Headline

Side	Passed	Rate	vs baseline
Baseline	12/36	33%	-
simplicio-cli (6-layer)	23/36	63%	+30 pts
simplicio-prompt (Yool)	8/36	22%	-11 pts

Per-model (pass = full PHPUnit suite green)

Model	Baseline	cli	sp	D cli	D sp
meta-llama/Llama-3.2-1B-Instruct	0/4 (0%)	0/4 (0%)	0/4 (0%)	+0	+0
google/gemma-3n-E4B-it	0/4 (0%)	0/4 (0%)	0/4 (0%)	+0	+0
google/gemma-3-4b-it	0/4 (0%)	3/4 (75%)	0/4 (0%)	+75	+0
qwen/qwen-2.5-7b-instruct	0/4 (0%)	1/4 (25%)	1/4 (25%)	+25	+25
meta-llama/Llama-3.1-8B-Instruct	2/4 (50%)	4/4 (100%)	0/4 (0%)	+50	-50
google/gemma-3-12b-it	2/4 (50%)	3/4 (75%)	1/4 (25%)	+25	-25
google/gemini-3.5-flash	3/4 (75%)	4/4 (100%)	1/4 (25%)	+25	-50
anthropic/claude-opus-4.7	2/4 (50%)	4/4 (100%)	2/4 (50%)	+50	+0
openai/gpt-5.5	3/4 (75%)	4/4 (100%)	3/4 (75%)	+25	+0

Per-task x model (base / cli / sp)

Task	Llama-3.2	gemma-3n	gemma-3	qwen-2.5	Llama-3.1	gemma-3	gemini-3	claude-op	gpt-5.5
password_strength	././.	././.	./P/.	./P/.	P/P/.	./P/.	P/P/.	./P/.	P/P/P
password_require_symbol	././.	././.	././.	././P	P/P/.	P/./P	P/P/P	./P/.	P/P/P
env_get_int	././.	././.	./P/.	././.	./P/.	./P/.	./P/.	P/P/P	./P/.
env_get_bool	././.	././.	./P/.	././.	./P/.	P/P/.	P/P/.	P/P/P	P/P/P