

Execution benchmark - real project, real tasks, real test suite

Date: 2026-05-28 Target: wesleysimplicio/sistema-sindico (real PHP 8 condominium system, public on GitHub) Tasks: 4 additive engineering changes across Core / Middleware / Repository / Router.

Methodology - what 'pass' actually means

NOT regex on model output. NOT a toy harness in isolation. The model's generated file replaces the real source file in a working copy of the real repo; a hidden PHPUnit test (never shown to the model, asserting TRUE and FALSE states) is added to tests/unit/Core/Hidden/; the FULL production suite runs (existing tests + hidden). Pass = vendor/bin/phpunit exit code 0 - same green/red the project's CI would use to merge a PR. The model's change must be correct AND must not break any pre-existing test of the real codebase.

All sides emit the complete file (identical output shape). Variables: baseline = raw goal; cli = the simplicio-cli 6-layer task contract.

Headline

Side	Passed	Rate	vs baseline
Baseline	12/36	33%	-
simplicio-cli (6-layer)	23/36	63%	+30 pts

Per-model (pass = full PHPUnit suite green)

Model	Baseline	simplicio-cli	Delta(pts)
meta-llama/Llama-3.2-1B-Instruct	0/4 (0%)	0/4 (0%)	+0
google/gemma-3n-E4B-it	0/4 (0%)	0/4 (0%)	+0
google/gemma-3-4b-it	0/4 (0%)	3/4 (75%)	+75
qwen/qwen-2.5-7b-instruct	0/4 (0%)	1/4 (25%)	+25
meta-llama/Llama-3.1-8B-Instruct	2/4 (50%)	4/4 (100%)	+50
google/gemma-3-12b-it	2/4 (50%)	3/4 (75%)	+25
google/gemini-3.5-flash	3/4 (75%)	4/4 (100%)	+25
anthropic/claude-opus-4.7	2/4 (50%)	4/4 (100%)	+50
openai/gpt-5.5	3/4 (75%)	4/4 (100%)	+25

Per-task x model (base / cli)

Task	Llama-3.2	gemma-3n	gemma-3	qwen-2.5	Llama-3.1	gemma-3	gemini-3	claude-op	gpt-5.5
password_strength	./.	./.	./P	./P	P/P	./P	P/P	./P	P/P
password_require_symbol	./.	./.	./.	./.	P/P	P/.	P/P	./P	P/P
env_get_int	./.	./.	./P	./.	./P	./P	./P	P/P	./P
env_get_bool	./.	./.	./P	./.	./P	P/P	P/P	P/P	P/P