

simplicio bench v13 — 5-side comparison (baseline / cli / cli+sp / cli+ag / cli+sp+ag)

Date: 2026-05-30

Sides

sem = baseline (raw goal). **com** = cli alone (6-layer). **sp** = sp runtime wrapping cli (composition). **ag** = cli + verify-loop (up to 3 attempts with PHPUnit/missing-pattern feedback). **spag** = sp-wrapped cli as the verify-loop seed (full stack composition + retry).

Models tested

- meta-llama/llama-3.2-3b-instruct (n=1 cases)

Headline — pct per side, both metrics

Model	Metric	base	cli	cli+sp	cli+ag	cli+sp+ag	Δcli	Δsp	Δag	Δsp+ag
llama-3.2-3b-instruct	exec	0%	0%	0%	0%	0%	+0	+0	+0	+0
llama-3.2-3b-instruct	regex	40%	80%	80%	80%	80%	+40	+40	+40	+40

Verify-loop convergence (avg attempts)

Model	cli+ag avg	cli+sp+ag avg
llama-3.2-3b-instruct	2.00	2.00

Cost & latency (exec, per call avg)

Model	Side	tokens/call	ms/call
llama-3.2-3b-instruct	baseline	950	2310
llama-3.2-3b-instruct	cli alone	1029	1729
llama-3.2-3b-instruct	cli + sp	1688	780
llama-3.2-3b-instruct	cli + ag	1982	3710
llama-3.2-3b-instruct	cli + sp + ag	2225	2197

Per-task × model (exec)

Format: **base/cli/cli+sp/cli+ag/cli+sp+ag** — P=pass, .=fail.

Task	llama-3.2-3b-instruct
password_strength	./././././.